075 Αρχιτεκτονική Υπολογιστών

Εισαγωγή στο προσομοιωτή Gem5

Ν. ΤΑΜΠΟΥΡΑΤΖΗΣ

ΤΜΗΜΑ ΗΛΕΚΤΡΟΛΟΓΩΝ ΜΗΧΑΝΙΚΩΝ & ΜΗΧΑΝΙΚΩΝ ΥΠΟΛΟΓΙΣΤΩΝ
ΑΡΙΣΤΟΤΕΛΕΙΟ ΠΑΝΕΠΙΣΤΗΜΙΟ ΘΕΣΣΑΛΟΝΙΚΗΣ

Introduction

- A full-system computer architecture simulator
  - Open source tool focused on architectural modeling
  - BSD license

- Encompasses
  - system-level architecture, as well as
  - processor micro-architecture.

- The gem5 simulation infrastructure is the merger of
  - The best aspects of the M5 and
  - The best aspects of GEMS

- M5
  - Highly configurable **simulation framework** to support multiple ISAs, and diverse CPU models
  - developed @ The University of Michigan

- GEMS [General Execution-driven Multiprocessor Simulator]
  - detailed and flexible **memory system model**
  - Includes support for multiple cache coherence protocols and interconnect models
  - developed @ The University of Wisconsin Madison

Where did it come from

Users and contributors

- Widely used in academia and industry

- Contributions from
  - ARM, AMD, Google,…
  - Wisconsin, Cambridge, Michigan, BSC, …

In a Nutshell, gem5...

… has had 11,558 commits made by 193 contributors
representing 386,321 lines of code

… is mostly written in C++
with a well-commented source code

… has a well established, mature codebase
maintained by a very large development team
with stable Y-O-Y commits

… took an estimated 104 years of effort (COCOMO model)
starting with its first commit in October, 2003
ending with its most recent commit 14 days ago

Publications with gem5

Languages

| C++ | 74% | Python | 18% |
| 13 Other | 8% | | |

Lines of Code

Code   Comments   Blanks

ARM

Why gem5?

- Runs real workloads
  - Analyze workloads that customers use and care about
  - … including complex workloads such as Android
- Comprehensive model library
  - Memory and I/O devices
  - Full OS
  - Clients and servers
- Rapid *early* prototyping
  - New ideas can be tested quickly
  - System-level impact can be quantified
- System-level insights
  - Enables us to study complex memory-system interactions
- Can be wired to custom models
  - *Add detail where it matters, when it matters!*

Ubuntu (Linux 4.x)

Android Nougat

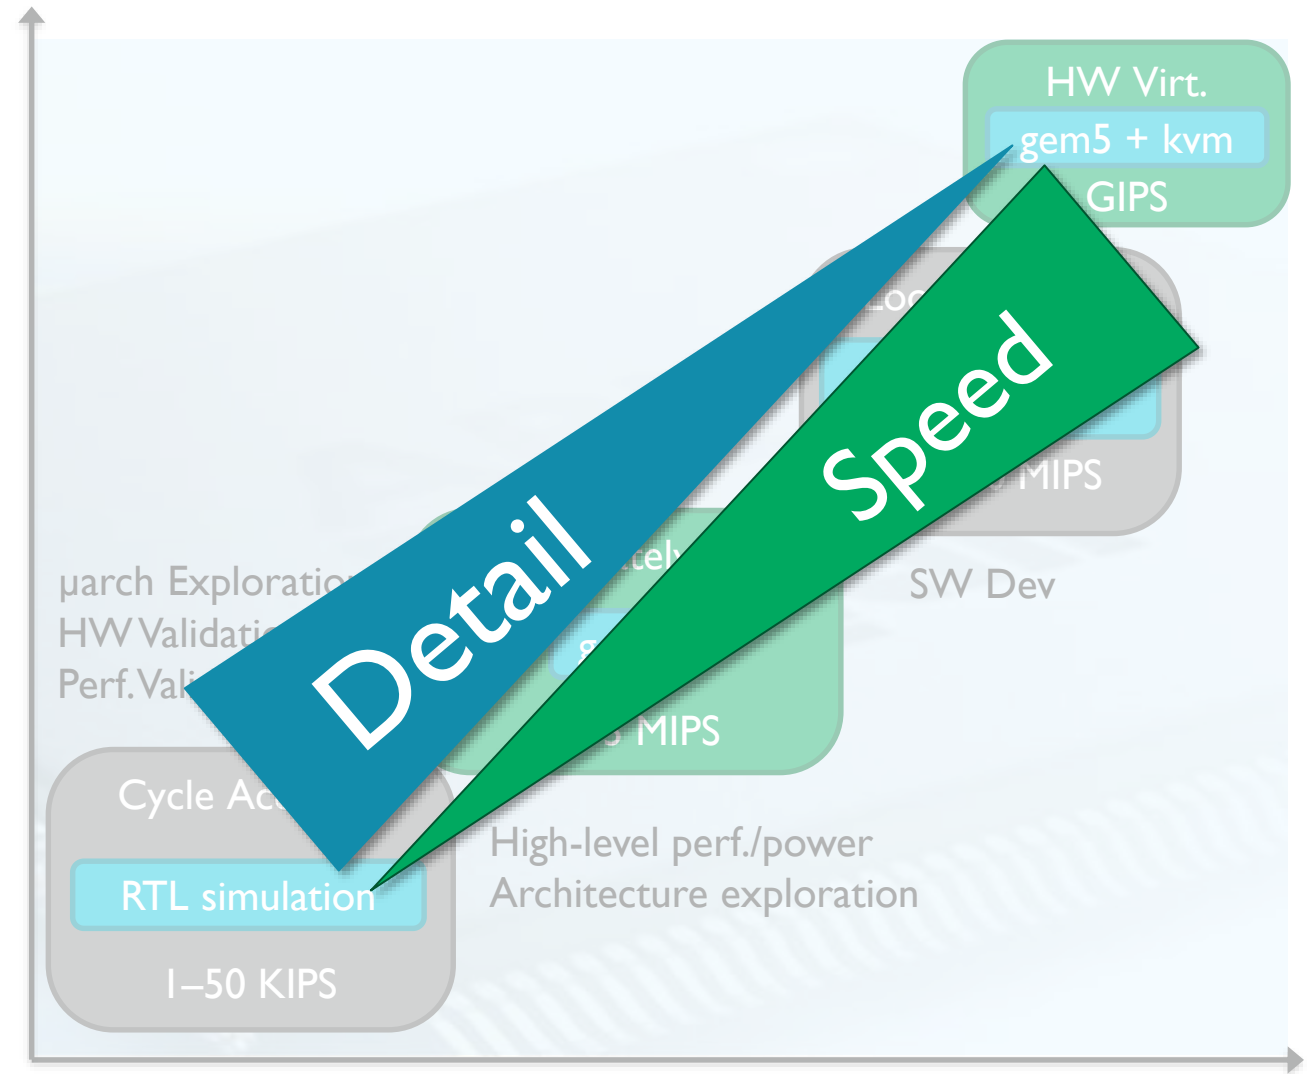# Level of detail

- **HW Virtualization**
  - Very no/limited timing
  - The same Host/guest ISA
- **Functional mode**
  - No timing, chain basic blocks of instructions
  - Can add cache models for warming
- **Timing mode**
  - Single time for execute and memory lookup
  - Advanced on bundle
- **Detailed mode**
  - Full out-of-order, in-order CPU models
  - Hit-under-miss, reodering, …

ARM

Prerequisites

- Operating system:
  - OSX, Linux
  - Limited support for Windows 10 with a Linux environment
- Software:
  - git
  - Python 2.7 (dev packages)
  - SCons
  - gcc 4.8 or clang 3.1 (or newer)
  - SWIG 2.0.4 or newer
  - make
- Optional:
  - dtc (to compile device trees)
  - ARMv8 cross compilers (to compile workloads)
  - python-pydot (to generate system diagrams)

ARM

Compiling gem5

```
$ scons build/ARM/gem5.opt -j4
```

- Guest architecture
- Several architectures in the source tree.
- Most common ones are:
  - **ARM**
  - **X86** – Popular in academia, but *very* strange timing behavior

- Optimization level:
  - **debug**: Debug symbols, no/few optimizations
  - **opt**: Debug symbols + most optimizations
  - **fast**: No symbols + even more optimizations

ARM

Modes

- gem5 has two fundamental modes
- Full system (FS)
  - For booting operating systems
  - Including devices
  - Interrupts, exceptions, privileged instructions
- Syscall emulation (SE)
  - For running individual applications, or set of applications on MP
  - Models user-visible ISA plus common system calls
  - System calls emulated, typ. by calling host OS
  - Simplified address translation model, no scheduling

Capabilities

- **Execution modes**: System-call Emulation (SE) & Full-System (FS)
- **ISAs**: Alpha, ARM, MIPS, Power, SPARC, x86
- **CPU models**: AtomicSimple, TimingSimple, InOrder, and O3
- **Cache coherence protocols**: broadcast-based, directories, etc.
- **Interconnection networks**: Simple & Garnet (Princeton, MIT)
- **Devices**: NICs, IDE controller, etc.
- **Multiple systems**: communicate over TCP/IP

Flexibility

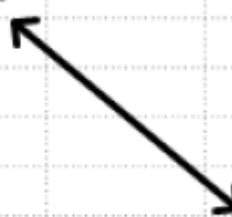| Processor		Memory System		
CPU Model	System Mode	Classic	Ruby	
			Simple	Garnet
Atomic Simple	SE	Speed		
	FS			
Timing Simple	SE			
	FS			
In-Order	SE			
	FS			
O3	SE			Accuracy
	FS			

Figure 1: Speed vs. Accuracy Spectrum.

- Classic (from M5): Fast and configurable memory system model.
- Ruby (from GEMS) : framework/infrastructure to model variety of cache-coherent memory system.

Example disk images

- Example kernels and disk images can be downloaded from gem5.org/Download
  - This includes pre-compiled boot loaders
  - Old but useful to get started
- Download and extract this into a new directory:
  - `wget http://www.gem5.org/dist/current/arm/aarch-system-2014-10.tar.xz`
  - `mkdir dist; cd dist`
  - `tar xvf ../aarch-system-2014-10.tar.xz`
- Set the M5_PATH variable to point to this directory:
  - `export M5_PATH=/path/to/dist`
- Most example scripts try to find files using `M5_PATH`
  - Kernels/boot loaders/device trees in `${M5_PATH}/binaries`
  - Disk images in `${M5_PATH}/disks`

ARM

Compiling Linux for gem5

```
1.  sudo apt install gcc-aarch64-linux-gnu
2.  git clone -b gem5/v4.4  https://github.com/gem5/linux-arm-gem5
3.  cd linux-arm-gem5
4.  make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- gem5_defconfig
5.  make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- -j `nproc`
```

- Builds the default kernel configuration for gem5
  - Has support for most of the devices that gem5 supports

ARM

# Running an example script

```
$ build/ARM/gem5.opt configs/example/arm/fs_bigLITTLE.py \
    --kernel path/to/vmlinux \
    --cpu-type atomic \
    --dtb $PWD/system/arm/dt/armv8_gem5_v1_big_little_1_1.dtb \
    --disk your_disk_image.img
```

- Simulates a bL system with 1+1 cores
  - Uses a functional 'atomic' CPU model
  - Use the 'timing' CPU type for an example OoO + InO configuration

ARM

Design philosophy

- gem5 is conceptually a Python library implemented in C++
    - **Configured** by instantiating Python classes with matching C++ classes
    - **Model parameters** exposed as attributes in Python
    - **Running** is controlled from Python, but implemented in C++

- Configuration and running are two distinct steps
    - **Configuration phase** ends with a call to instantiate the C++ world
    - **Parameters cannot be changed** after the C++ world has been created

ARM

Control flow
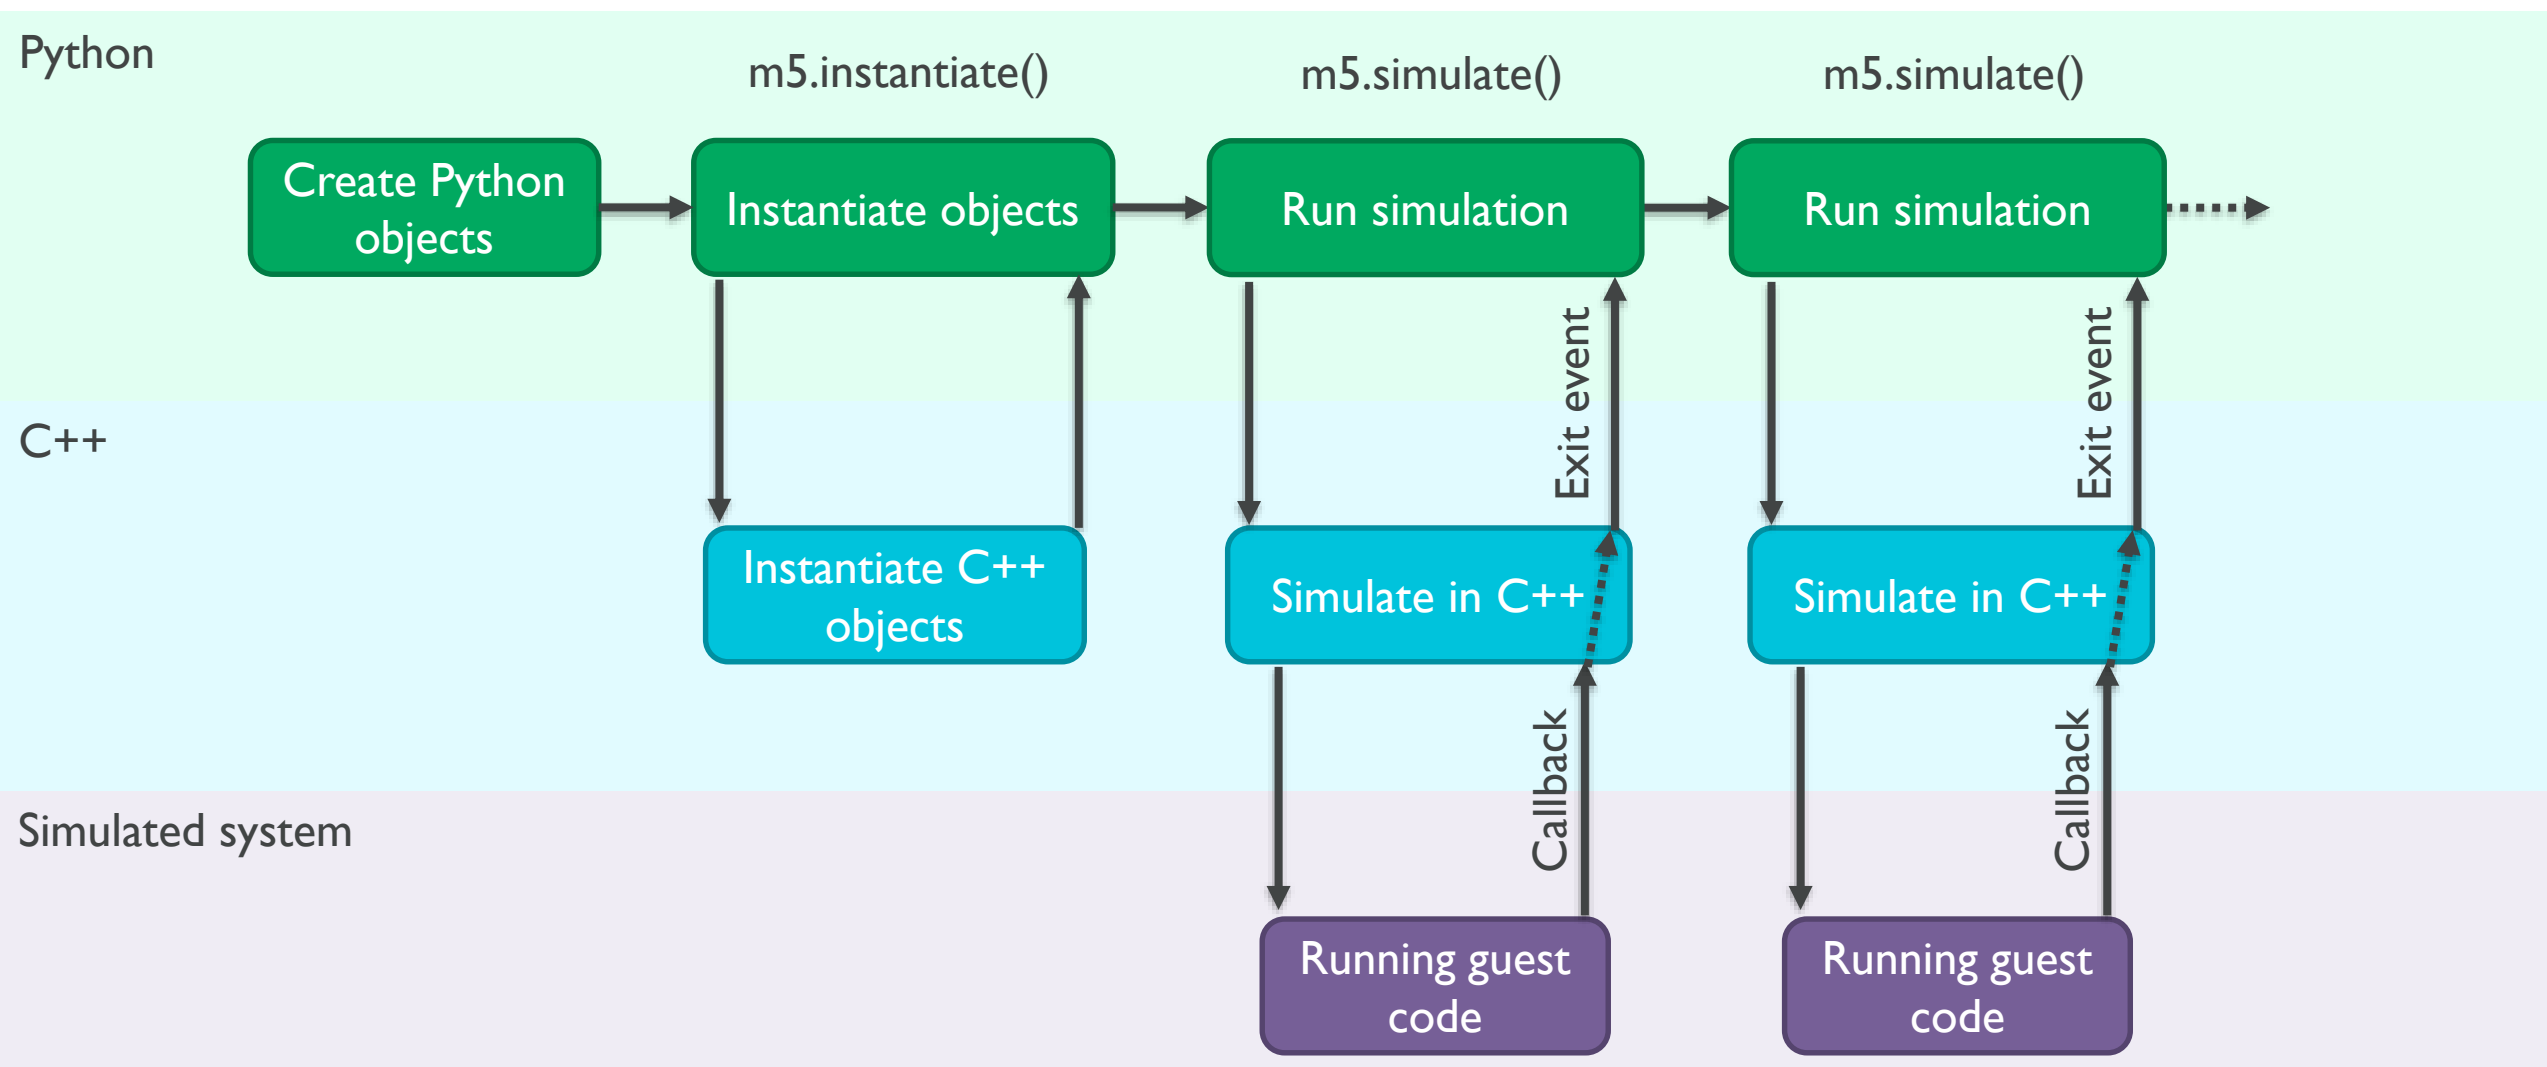
Overriding model parameters

import m5

```
class L1DCache(m5.objects.Cache):
    assoc = 2
    size = '16kB'
```

- Use gem5's base Cache
- Override associativity
- Override size

```
class L1ICache(L1DCache):
    assoc = 16
```

- Use defaults from L1DCache
- Override associativity again

```
l1i = L1ICache(assoc=8,
        repl=m5.objects.RandomRepl())
```

- Override parameters at instantiation time
- We'll cover memory ports later

ARM

Running

```
m5.instantiate()
```

- Instantiate the C++ world

```
event = m5.simulate()
```

- Start the simulation

```
print 'Exiting @ tick %i: %s' \
    % ( m5.curTick(),
        event.getCause())
```

- Print why the simulator exited
- Sometimes desirable to call m5.simulate() again.

```
m5.simulate(m5.tick.fromSeconds(0.1))
```

- Run for a fixed number of simulated seconds

ARM

Creating Checkpoints

```
m5.checkpoint('name.cpt')
```

- Checkpoints can be used to **store the simulator's state**
  - Can be used to implement SimPoints or similar methodologies

- Checkpoint limitations:
  - The act of taking a checkpoint affects system state!
  - **Checkpoints don't store cache state**
  - **Checkpoints don't store pipeline state**

ARM

Restoring Checkpoints

```
m5.instantiate('name.cpt')
```

- Instantiate system and load state from checkpoint

```
event = m5.simulate()
```

- Run in the same way as before

ARM

Dumping statistics

- Can be requested from Python:
  - `m5.stats.dump()`: Dump statistics
  - `m5.stats.reset()`: Reset stat counters

- Guest command line:
  - `m5 dumpstats [[delay] [period]]`
  - `m5 dumpresetstas [[delay] [period]]`

- Guest code using libm5.a:
  - `m5_dump_stats(delay, periodicity)`: Dump statistics
  - `m5_dumpreset_stats(delay, periodicity)`: Dump & reset statistics

ARM

Understanding gem5 output

```
> ls m5out

config.ini    config.json    stats.txt
```

config.ini: Dumps all of the parameters of all SimObjects (modules). This shows exactly what you simulated.

config.json: Same as config.ini, but in json format.

stats.txt: Detailed statistic output. Each SimObject defines and updates statistics. They are printed here at the end of simulation.

stats.txt

```
---------- Begin Simulation Statistics ----------
sim_seconds        0.000346        # Number of seconds simulated
sim_ticks          345518000       # Number of ticks simulated
final_tick         345518000       # Number of ticks from
sim_freq           1000000000000   # Frequency of simulate
...
sim_insts          5712            # Number of instructions simulated
sim_ops            10314           # Number of ops (including micro ...
...
system.mem_ctrl.bytes_read::cpu.inst  58264   # N...
system.mem_ctrl.bytes_read::cpu.data  7167    # N...
...
system.cpu.committedOps            10314   # Number o...
system.cpu.num_int_alu_accesses 10205  # Number of integer ...
```

sim_seconds: name of stat. This shows *simulated guest* time

Every SimObject can have its own stats. Names are what you used in the Python config file

Debugging

Debugging Facilities

- Tracing
  - Instruction tracing
  - Diffing traces

- Using gdb to debug gem5
  - Debugging C++ and gdb-callable functions
  - Remote debugging

- Pipeline viewer

ARM

Tracing/Debugging

- `printf()` is a nice debugging tool
  - Keep good print statements in code and selectively enable them
  - Lots of debug output can be a very good thing when a problem arises
  - Use `DPRINTF`s in code
  - `DPRINTF(TLB, "Inserting entry into TLB with pfn:%#x…)`

- Example flags:
  - `Fetch, Decode, Ethernet, Exec, TLB, DMA, Bus, Cache, O3CPUAll`
  - Print out all flags with `./build/ARM/gem5.opt -- debug-help`

- Enabled on the command line
  - `--debug-flags=Exec`
  - `--debug-start=30000`
  - `--debug-file=my_trace.out`
  - Enable the flag Exec; Start at tick `30000`; Write to `my_trace.out`

ARM

Sample Run with Debugging

Command Line:

```
22:44:28 [/work/gem5] ./build/ARM/gem5.opt --debug-flags=Decode --debug-start=50000-- debug-file=my_trace.out configs/example/se.py  -c tests/test-progs/hello/bin/arm/linux/hello
…
**** REAL SIMULATION ****
info: Entering event queue @ 0.  Starting simulation...
Hello world!
Exiting @ tick 3107500 because target called exit()
```

my_trace.out:

```
2:44:47 [ /work/gem5] head m5out/my_trace.out
  50000:      system.cpu:  Decode:  Decoded cmps instruction:     0xe353001e
  50500:      system.cpu:  Decode:  Decoded ldr instruction: 0x979ff103
  51000:      system.cpu:  Decode:  Decoded ldr instruction: 0xe5107004
  51500:      system.cpu:  Decode:  Decoded ldr instruction: 0xe4903008
  52000:      system.cpu:  Decode:  Decoded addi_uop instruction: 0xe4903008
  52500:      system.cpu:  Decode:  Decoded cmps instruction:     0xe3530000
  53000:      system.cpu:  Decode:  Decoded b instruction:    0x1affff84
  53500:      system.cpu:  Decode:  Decoded sub instruction: 0xe2433003
  54000:      system.cpu:  Decode:  Decoded cmps instruction:     0xe353001e
  54500:      system.cpu:  Decode:  Decoded ldr instruction: 0x979ff103
```

ARM

Adding Your Own Flag

- ## Print statements put in source code

    - Encourage you to add ones to your models or contribute ones you find particularly useful

- ## Macros remove them from the `gem5.fast` binary

    - There is no performance penalty for adding them
    - To enable them you need to run `gem5.opt` or `gem5.debug`

- ## Adding one with an existing flag

    - `DPRINTF(<flag>, "normal printf %s\n", "arguments");`

- ## To add a new flag add the following in a `Sconscript`

    - `DebugFlag('MyNewFlag')`
    - `Include corresponding header, e.g. #include "debug/MyNewFlag.hh"`

ARM